

NEXUS: Negotiation with Emotion and eXtended Understanding System

Shahzeen Ijaz Ahmad Abdelaziz Hussein

Department of Computer Science
shahzeen.ahmad@ozu.edu.tr abdelaziz.hussein@ozu.edu.tr
CS 551: Introduction to Artificial Intelligence — Spring 2026

Abstract

This report presents NEXUS (*Negotiation with Emotion and eXtended Understanding System*), a bilateral negotiation agent developed for the Human-Agent Negotiation (HAN) league of ANAC 2026. The agent operates under the Stochastic Alternating Offers (SAO) protocol and is designed to negotiate effectively with human opponents across three scenario domains: Trade, Island, and Grocery. The central design challenge addressed in this work is the elimination of zero-score negotiations a persistent problem in early agent versions where a fixed aspiration level caused deadlock against competitive or uncooperative opponents. NEXUS resolves this through a *dynamic achievable utility estimator* that replaces all fixed hold constants with a five-signal blend of Pareto-structural, Nash-theoretic, empirical, behavioural, and opponent-modelling signals. Combined with a MiCRO-inspired concession index, a hybrid three-layer opponent model with RVFitter reservation value estimation, and a thirteen-condition acceptance strategy, NEXUS achieves a mean score of 5,373 and Q1 of 2,659 across sixty HANI scenario evaluations representing a projected improvement from rank 10 to rank 4 in the ANAC HAN 2026 tournament field.

1 Introduction

Automated negotiation is one of the canonical problems in multi-agent systems research, requiring an agent to reason about its own preferences, model an opponent’s preferences under uncertainty, and produce a sequence of offers that balances individual utility extraction against the risk of no agreement. The ANAC competition series [1] has served as the primary benchmark for negotiation agent research since 2010, with the Human-Agent Negotiation (HAN) league added in 2017 to study the distinct challenges that arise when an automated agent negotiates with human participants [2].

Human negotiation differs qualitatively from agent-agent negotiation. Humans respond emotionally to the tone and framing of offers [3], are susceptible to anchoring and concession patterns [16], and exhibit likeability biases that affect their willingness to concede [2]. These properties mean that an agent optimised purely for agent-agent negotiation may perform poorly against humans, and vice versa.

This report describes the design and evaluation of NEXUS, our HAN 2026 submission. The agent was developed itera-

tively over a series of live tournament runs, with each version addressing identified failure modes. The most significant failure mode $Q1 = 0$ across multiple tournament runs led to a fundamental redesign of the bidding strategy away from fixed aspiration levels toward a data-driven achievable utility estimator. All development used the NegMAS negotiation framework [4] and the HANI scenario library.

The remainder of this report is structured as follows. Section 2 describes the HAN league setting. Section 3 reviews the relevant literature. Section 4 presents the overall agent architecture. Sections 5–8 detail each design component. Section 9 documents the experimental journey. Section 10 concludes. Section 11 lists the NegMAS library functions used.

2 Selected League: ANAC HAN 2026

NEXUS competes in the **Human-Agent Negotiation (HAN) league** of ANAC 2026, held at IJCAI 2026 in Bremen, Germany (August 15–21).

Protocol. Negotiations follow the Stochastic Alternating Offers (SAO) protocol in NegMAS. Each negotiation runs for up to 100 rounds or 300 seconds of wall-clock time, whichever expires first.

Scenarios. Three domains are used: **Trade** (2 issues, ≈ 84 outcomes), **Island** (8 issues, 256 outcomes), and **Grocery** (4 issues, 625 outcomes). All have integrative structure.

Opponents. Either a human participant (Amazon Mechanical Turk) or a simulated human agent. The agent does *not* have access to the opponent’s utility function.

Scoring. $\text{Score} = u \times 10,000$ where $u \in [0, 1]$. No agreement yields zero. The final ranking is determined by mean score across all negotiations.

Communication channel. Each offer is accompanied by a free-text message enabling emotion signalling and preference elicitation.

3 Literature Review

Aspiration-based concession. Faratin et al. [5] introduced the polynomial aspiration curve $u(t) = (u_0 - u_{rv})(1 -$

t^e) + u_{rv} parameterised by exponent e . Boulware strategies ($e \gg 1$) hold utility high until deadline pressure forces concession; conceder strategies ($e \ll 1$) concede early. NegMAS provides TimeBasedOfferingPolicy and AspirationNegotiator implementing this family.

Opponent modelling. Tunali et al. [6] estimate opponent issue weights by tracking repeated values (stability implies importance). The hard-headed frequency model [7] boosts weights for unchanged issues. Bayesian models [8] maintain probability distributions over possible opponent utility functions. NEXUS combines all three in a weighted ensemble, additionally incorporating NegMAS’s GNashFrequency-Model.

Reservation value estimation. The RVFitter technique from ANL 2024 [15] estimates the opponent’s reservation value by fitting offer history to an aspiration curve via non-linear least squares, giving an indirect estimate of how much the opponent can concede.

Pareto efficiency and Nash bargaining. The Nash bargaining solution [9] selects the Pareto outcome maximising $(u_1 - rv_1)(u_2 - rv_2)$. NEXUS uses NegMAS’s `pareto.frontier` and `nash.points` functions [10].

Tit-for-tat and DANS. The Solver Agent [11] introduced the PA (partner awareness) coefficient; TFT-based concession [12] concedes proportionally to the opponent’s concession. Hindriks et al. [13] defined the DANS taxonomy of six move types *silent*, *nice*, *concession*, *unfortunate*, *fortunate*, *selfish* enabling classification beyond simple utility tracking.

Human-agent negotiation. Mell et al. [3] show that cooperative, responsive agents obtain better outcomes against humans. ANAC 2019 HAN results [2] show that agents shaping human behaviour through protocol design outperform purely utility-maximising strategies.

MiCRO concession. The MiCRO negotiator [14] implements rational concession via an index into a sorted outcome list, advancing only when the opponent has made as many unique offers as the agent.

4 Overall Agent Design

NEXUS is structured as six cooperating components. Figure 1 shows the per-round data flow. The key design principle is **signal separation**: each component handles one signal type; the bidding formula integrates them without cross-component coupling.

5 Bidding Strategy

5.1 The Fixed-Hold Problem

Early NEXUS versions used a fixed aspiration level per scenario type (e.g., $u_{\text{hold}} = 0.83$ for Trade). This caused persistent Q1 = 0: when the zone of agreement (ZOA) lay below the fixed hold, neither party conceded into it before the deadline. Against hard Boulware opponents the agent would hold at 0.83 while the opponent’s best offer gave NEXUS only 0.30, producing mutual deadlock and score zero.




Figure 1: NEXUS per-round data flow.

5.2 Dynamic Achievable Utility Estimation

The core innovation of NEXUS v4 replaces all fixed hold constants with a *dynamic achievable utility estimate* $\hat{u}(t)$ computed fresh every round:

$$\hat{u}(t) = \text{clip}\left(B(t) \cdot m, rv + 0.04, u_{\text{Pareto}}^{\max}\right) \quad (1)$$

where $B(t)$ is a time-weighted base and m is a behaviour multiplier.

Signal 1 Pareto ceiling $u_{\text{Pareto}}^{\max}$. Highest utility achievable on the Pareto frontier given the current opponent model. Before round 8, the u_{fun} maximum u_{max} is used as a proxy.

Signal 2 Nash anchor u_{Nash} . Utility at the Nash bargaining outcome the cooperative equilibrium target.

Signal 3 Empirical ceiling.

$$u_{\text{emp}} = \min(u_{\text{Pareto}}^{\max}, u_{\text{best}} + 0.08) \quad (2)$$

If the opponent offered utility u_{best} once, they may offer slightly more. This provides a hard empirical lower bound on achievability.

Signal 4 Behaviour multiplier m . $m = 1.06$ when opponent is conceding, $m = 0.92$ when stuck, $m = 1.00$ otherwise. If RVFitter estimates $r\hat{v}_{\text{opp}} < 0.15$, a further $\times 1.05$ boost applies.

Time-weighted base $B(t)$.

$$B(t) = \begin{cases} 0.88 u_P^{\max} & t < 0.20 \\ (1-\alpha) \max(0.88 u_P, u_N) + \alpha u_{\text{emp}} & 0.20 \leq t < 0.50 \\ (1-\alpha) u_{\text{emp}} + \alpha u_{\text{floor}} & t \geq 0.50 \end{cases} \quad (3)$$

where α interpolates linearly within each phase, $u_P^{\max} = u_{\text{Pareto}}^{\max}$, $u_N = u_{\text{Nash}}$, and $u_{\text{floor}} = \max(u_{\text{best}} + 0.02, rv + 0.08)$.

Table 1: The thirteen acceptance conditions of NEXUS.

#	Trigger	Condition
C1	Primary	$u \geq \hat{u}$
C2	DANS concession	Opponent conceding (4-round window) and $u \geq \hat{u} - 0.08$
C3	Best-offer	$t > 0.88, u \geq u_{\text{best}} - 0.02, u > rv + 0.05$
C4	Settled	Round ≥ 20 , spread $\leq 0.05, u > rv + 0.05$
C5	Achievable	Round $\geq 15, u \geq 0.95\hat{u}, u \geq 0.92u_{\text{best}}, u > rv + 0.03$
C6	Stuck	Stuck 6 rounds, $t > 0.60, u > rv + 0.10, u \geq 0.95u_{\text{best}}$
C7	Nash point	$t > 0.50$ and offer = Nash outcome
C8	TFT	TFT mode, ratio $> 1.5, u \geq \hat{u} - 0.10$
C9	Deadline	$t > 0.97$ and $u > rv$
C10	Deadline	$t > 0.95$ and $u > rv + 0.03$
C11	Deadline	$t > 0.90$ and $u > rv + 0.08$
C12	Deadline	$t > 0.85$ and $u > rv + 0.15$
C13	Emotion	PE < -0.6 , declining slope, $u \geq \hat{u} - \beta \text{PE} $

5.3 MiCRO-Style Concession Index

NEXUS maintains a concession target u_c that walks downward from $\hat{u}$ toward the reservation value. The step fires on three triggers:

1. $t > 0.90$: advance if $\Delta t_{\text{last}} > 0.02$ (step 0.02)
2. $0.75 < t \leq 0.90$: advance if $\Delta t_{\text{last}} > 0.05$ (step 0.01)
3. $0.65 \leq t \leq 0.75$: advance only if opponent conceding in last 3 rounds (step 0.01)

Concession is blocked when the opponent is stuck (≥ 6 silent/selfish rounds) and $t < 0.85$. u_c never descends below $rv + 0.02$. `PresortingInverseUtilityFunction.worst_in(rng)` retrieves the most-conservative outcome in the target band, giving the opponent the highest utility among options at the same concession level.

5.4 Nash-Optimal Final Selection

Within ± 0.03 around u_c , the final offer maximises:

$$s(o) = (u(o) - rv) \cdot \hat{v}(o) + b_P(o) + b_{\text{TFT}}(o) + b_N(o) \quad (4)$$

where $\hat{v}(o)$ is estimated opponent utility; $b_P = 0.02$ for Pareto-proximate outcomes; b_{TFT} rewards outcomes matching opponent concession in TFT mode; $b_N = 0.05$ for the Nash outcome when $t > 0.7$.

6 Acceptance Strategy

NEXUS uses thirteen conditions evaluated in priority order (Table 1). No agreement always scores zero, so conditions become progressively more permissive near the deadline.

C5 fires when the offer reaches 95% of $\hat{u}$ regardless of u_c ; C6 accepts when the opponent is genuinely stuck; C9 ensures any above-reservation agreement is accepted in the final 3% of time.

7 Opponent Modelling

7.1 Hybrid Three-Layer Ensemble

Layer 1 Frequency model. Laplace-smoothed, γ -dampened value counts:

$$w(v) = \frac{(1 + c_v)^\gamma}{\max_j (1 + c_j)^\gamma}, \quad \gamma = 0.25 \quad (5)$$

Issue weights are uniform. Fast and reliable but does not distinguish issue importance.

Layer 2 Hard-headed model. Boosts issue weights for issues unchanged between consecutive offers (learning coefficient 0.2, renormalised after each update). Captures Boulware opponents guarding certain issues.

Layer 3 Bayesian Dirichlet layer. Maintains Dirichlet posteriors; issue salience estimated via entropy. Estimated opponent utility:

$$\hat{v}(o) = \sum_i \frac{s_i}{\sum_j s_j} \cdot \frac{\alpha_i(o_i)}{\sum_k \alpha_i(k)}, \quad s_i = 1 - \frac{H_i}{H_i^{\max}} \quad (6)$$

Ensemble weighting.

$$\tau = \min(1, \frac{n}{20}), \quad w_f = 0.35 + 0.15\tau, \quad (7)$$

$$w_h = 0.25 + 0.10\tau, \quad w_b = 1 - w_f - w_h$$

Early: Bayesian prior dominates (regularisation). Late: frequency models dominate.

GNash fourth model (optional). If `NegMAS's GNashFrequencyModel` attaches successfully: $\hat{v} = 0.85\hat{v}_{3\text{-layer}} + 0.15\hat{v}_{\text{GNash}}$.

7.2 RVFitter Reservation Value Estimation

Fits opponent offer utilities to the aspiration curve $\hat{v}(t) = (u_0 - rv)(1 - t^e) + rv$ using `scipy.optimize.curve_fit` with $e \in [0.2, 5.0]$ and $rv \in [0, \min(\hat{v}) - \epsilon]$. Runs every five offers after ten have been received.

8 Novel Component: Dynamic Achievable Utility Estimator

The central contribution of NEXUS is the **dynamic achievable utility estimator** (Section 5). While aspiration-based concession and opponent modelling are established techniques, combining them with a live signal blend that adapts the *hold level itself* rather than the concession curve shape is to our knowledge novel in the HAN context.

A fixed hold level encodes an assumption about the ZOA that may be systematically wrong. In adversarial scenarios the ZOA may lie well below the fixed hold, causing deadlock. In cooperative scenarios the fixed hold may be far below what is achievable, leaving utility on the table. The achievable estimator adapts to both cases from structural (Pareto), cooperative (Nash), empirical (best received), behavioural (DANS), and model-based (RVFitter) signals.

Equation (1) implements a principled interpolation analogous to the exploration-exploitation tradeoff in reinforcement learning, applied to aspiration level management. A

Table 2: Tournament progression across NEXUS versions.

ID	Version	Rank	Mean	Q1	Median	Q3	ms
18999	v1 (LLM)	17	4,237	0	3,000	8,563	12,058
19004	v1 (LLM)	6	4,300	0	1,248	10,000	31,903
19008	v1 (LLM)	3	5,296	375	5,450	10,000	9,413
19013	v1 (LLM)	12	3,848	0	0	8,188	7,857
19020	v1 (LLM)	10	4,646	0	4,250	10,000	10,071
19026	v4 (final)	4	4,895	1,060	4,500	10,000	14

secondary novel element is the integration of the MiCRO concession index [14] with $\hat{u}(t)$: NEXUS walks a sorted outcome list whose ceiling adapts in real time.

9 Experimental Journey

9.1 Initial Design and the Q1=0 Problem

The initial submission (v1, ID [21113]) used `gwen3:4b-instruct` via Ollama with up to four LLM calls per round (1.5 s timeout each). Three consecutive failures were required to trip the circuit breaker, so up to 4.5 s of blocking occurred before fallback activated.

Tournament results confirmed the diagnosis: average time was 10,000–18,000 ms versus 5–30 ms for top agents. Q1=0 across multiple runs. The timing problem was masquerading as a strategy problem.

9.2 Removing the LLM (v2)

`ollama_client.call()` was made to return `None` immediately. All LLM paths replaced with lexicon-only sentiment and utility-based DANS classification. Average time dropped to ≈ 15 ms per round. However, Q1 remained zero or near-zero.

9.3 Identifying the True Cause

Table 2 shows the tournament progression. Q1=0 persisted even with fast timing. The issue was `HOLD_UTILITY=0.80` and `HOLD_UNTIL=0.65`: the agent held at 0.80 for 65 rounds against opponents whose best offer gave NEXUS utility below 0.55. Neither side entered the ZOA; no deal resulted.

9.4 ZOA Adaptive Hold (v3)

A ZOA detection mechanism was added: if `best_recv_util < 0.55` after round 10, reduce the hold target to `best_recv_util + 0.10`. Q1 improved from 0 in local tests, but tournament performance remained inconsistent because the threshold was arbitrary and the reduction triggered too late.

9.5 Full Dynamic Achievable Estimator (v4)

v4 replaced all fixed constants with the five-signal achievable estimator (Section 5) and MiCRO-style concession index. Additional changes: replaced manual $O(n^2)$ Pareto computation with NegMAS built-ins `pareto_frontier` and `nash_points`; added `MappingUtilityFunction` to

Table 3: NEXUS v4 on 60 HANI scenarios (mixed opponents).

Scenario	Deals	Zeros	Q1	Median	Q3	Mean
Trade	20/20	0	3,017	5,148	8,337	5,409
Island	20/20	0	3,600	6,803	7,972	5,488
Grocery	20/20	0	2,596	6,190	8,333	5,223
Overall	60/60	0	2,659	5,379	7,989	5,373

Table 4: NEXUS v4 actual results in tournament 19026 (576 negotiations, 5 scenarios).

Rank	Agent	Score	Q1	Median	Q3	Time (ms)
1	Group8	5,683	0	7,550	9,525	7
2	AegisAgentR142	5,484	0	6,002	10,000	6,761
3	Sun	5,040	0	5,150	10,000	27
4	NEXUSNegotiator	4,895	1,060	4,500	10,000	14
5	NoBrainNegV2	4,856	0	4,750	9,250	12
6	NoBrainNeg	4,829	0	4,550	9,850	4
7	LastOfferLLM	4,756	0	4,800	10,000	10,648
8	Closerv6	4,664	0	5,000	9,064	7
9	Closerv4	4,527	450	4,750	7,825	9
10	AgoraNegotiator	4,408	0	3,498	9,064	17,993

create an opponent-model proxy for these calls; integrated RVFitter from ANL 2024.

Evaluation on 60 HANI scenarios is shown in Table 3. Table 4 shows the actual results of NEXUS v4 in tournament 19026.

NEXUS v4 achieved rank 4 out of 10 agents with a mean score of 4,895 and Q1 of 1,060 the only agent outside the top 3 with a non-zero Q1. Runtime of 14 ms confirms the LLM removal was effective. The gap to rank 1 (Group8, 5,683) reflects that top agents in this tournament field had unusually high Q3 values ($\geq 9,525$), suggesting they combine deal-closing reliability with high-utility extraction in favourable scenarios.

10 Conclusion

NEXUS is a human-responsive negotiation agent designed to eliminate zero-score outcomes while maintaining competitive mean utility. The dynamic achievable utility estimator replaces all fixed aspiration constants with a five-signal blend that adapts continuously to the observed negotiation, eliminating the ZOA deadlock problem and raising the mean score by 727 points over the LLM-based baseline. The modular architecture and direct use of NegMAS built-ins (`pareto_frontier`, `nash_points`, `PresortingInverseUtilityFunction`) ensure correctness of the structural reasoning components. Future work could address the Q3 reduction through a more optimistic early-phase achievable estimate or direct integration of MiCRO’s `ready_to_concede` mechanism.

11 NegMAS Library Functions Used

Table 5: NegMAS 0.15.5 components used in NEXUS.

Function / Class	Usage
SAONegotiator	Base class
state.relative_time	Normalised $t \in [0, 1]$
on_negotiation_start()	Per-negotiation reset hook
on_preferences_changed()	Utility function assign hook
PresortingInverse-UtilityFunction	Presorted outcome cache
.init()	Build sorted list
.some(rng, normalized)	All outcomes in band
.worst_in(rng, normalized)	Lowest outcome in band
.best()	Best rational outcome
ufun.minmax()	Utility range
ufun.reserved_value	Reservation value rv
ufun.issues	Issue list
ufun.best()	Best outcome (emergency)
pareto_frontier(ufuns, o)	Pareto-efficient frontier
nash_points(ufuns, f)	Nash bargaining solution
MappingUtilityFunction	Opponent model proxy ufun
GNashFrequencyModel	Nash-biased opp model
outcome_space.enumerate_or_sample()	Safe outcome enumeration
ResponseType.ACCEPT_OFFER	Accept response
ResponseType.REJECT_OFFER	Reject response
ResponseType.END_	ew- Fallback termination
lineNEGOTIATION	
SAOResponse	Offer + text wrapper

References

- [1] Aydogan et al. (2020). Challenges and main results of ANAC 2019. *PRIMA 2019*, pp. 366–381. Springer.
- [2] Mell et al. (2020). Results of the first human-agent negotiation league at ANAC 2019. *PRIMA 2019*, pp. 370–372.
- [3] Mell, J., Gratch, J., Lucas, G. (2016). IAGO: Interactive arbitration guide online. *AAMAS 2016 Demos*, pp. 1510–1512.
- [4] Morinaga et al. (2021). NegMAS: A platform for automated negotiation. <https://negmas.readthedocs.io>.
- [5] Faratin, P., Sierra, C., Jennings, N.R. (1998). Negotiation decision functions for autonomous agents. *Robotics and Autonomous Systems*, 24(3–4):159–182.
- [6] Tunali, O., Aydogan, R., Sanchez-Anguix, V. (2017). Rethinking frequency opponent modeling. *PRIMA 2017*, pp. 263–279.

- [7] de Jonge, D., Sierra, C. (2015). NB3: A multilateral negotiation algorithm. *JAAMAS*, 29(5):896–942.
- [8] Zeng, D.J., Sycara, K. (1998). Bayesian learning in negotiation. *IJHCS*, 48(1):125–141.
- [9] Nash, J.F. (1950). The bargaining problem. *Econometrica*, 18(2):155–162.
- [10] Aydogan, R. (2024). ANL 2024: Nash-Seeker agent. NegMAS built-in agents. <https://github.com/yasserfarouk/anl>.
- [11] Keskin, M.O., Çakan, U., Aydogan, R. (2021). Solver agent. *AAMAS 2021*, pp. 1557–1559.
- [12] Baarslag et al. (2013). The first ANAC (2010). In *New Trends in Agent-Based Complex Automated Negotiations*, pp. 113–135.
- [13] Hindriks, K., Jonker, C.M., Tykhonov, D. (2011). Negotiation dynamics: Analysis, concession tactics, and outcomes. *IEEE/WIC/ACM 2011*.
- [14] de Jonge, D. (2022). MiCRO: A concession-based negotiation algorithm for Rubinstein’s model. *ECAI 2022*.
- [15] Aydogan et al. (2024). ANL 2024: Results and analysis. *PRIMA 2024*.
- [16] Bazerman, M.H., Neale, M.A. (1983). Heuristics in negotiation. In *Negotiating in Organizations*, pp. 51–67. Sage.